

EU AI Act Compliance Assessment

Risk Classification and Obligation Mapping

Version	1.0
Date	2026-05-06
Author	Jay Wingard
Subject System	Genesis Autonomous AI Agent
Classification	Public

Document History

Version / Date	Author	Description
1.0 (2026-05-06)	Jay Wingard	Initial release

Contents

1	EU AI Act Compliance Assessment: Genesis Autonomous AI Agent	2
1.1	1. Executive Summary	2
1.1.1	1.1.1 Important Caveat	2
1.2	2. Risk Classification Analysis	2
1.2.1	2.1 Prohibited Practices (Article 5)	2
1.2.2	2.2 High-Risk Classification (Article 6 and Annex III)	3
1.2.3	2.3 Limited Risk-Transparency Obligations (Article 50)	4
1.2.4	2.4 Risk Classification Summary	4
1.3	3. Obligation Mapping	5
1.3.1	3.1 Current Obligations (Limited Risk)	5
1.3.2	3.2 Hypothetical High-Risk Analysis (Articles 9-17)	6
1.3.3	3.3 Readiness Summary for High-Risk Obligations	6
1.3.4	3.4 Preparatory Remediation for Hypothetical High-Risk	7
1.4	4. General-Purpose AI Model Obligations (Chapter V)	7
1.4.1	4.1 Genesis and GPAI	7
1.4.2	4.2 Downstream Deployer Considerations	8
1.5	5. Territorial Scope	8
1.5.1	5.1 Applicability to Genesis	8
1.5.2	5.2 Practical Assessment	8
1.6	6. Timeline Analysis	9
1.7	7. The “Agent” Question	9
1.8	8. Limitations of This Assessment	10
1.9	9. Gap Summary and Remediation	10
1.9.1	9.1 Current Compliance Gaps	10
1.9.2	9.2 Remediation Plan	10
1.9.3	9.3 Cross-Reference to Other Assessments	11

EU AI Act Compliance Assessment: Genesis Autonomous AI Agent

Version: 1.0 | Date: 2026-05-06 | Classification: Public

1. Executive Summary

This document assesses the Genesis autonomous AI agent system against the EU Artificial Intelligence Act (Regulation (EU) 2024/1689), published in the Official Journal on 12 July 2024. The assessment covers risk classification, obligation mapping for the applicable risk tier, transparency requirements, and General-Purpose AI (GPAI) model obligations.

The EU AI Act creates a risk-based regulatory framework with four tiers: Prohibited, High-Risk, Limited Risk, and Minimal Risk. The Act addresses AI systems (products and services deployed within the EU or whose outputs are used within the EU) and GPAI models (foundation models used as components within AI systems).

Genesis presents an unusual case for EU AI Act classification. The system does not fit neatly into any Annex III high-risk use case category, and it does not meet the definition of a prohibited practice. Its most likely classification is Limited Risk under the transparency obligations, with potential high-risk implications depending on how its external-facing capabilities are used.

The system description (D0) provides the factual foundation for this assessment.

Important Caveat

The EU AI Act is new legislation with limited interpretive guidance. The Commission's implementing acts and harmonized standards are still in development. This assessment represents a good-faith analysis based on the text of the Regulation as published, the high-level summary from the Future of Life Institute, and general principles of EU technology regulation. It is not legal advice, and specific classifications may change as implementing guidance is published.

2. Risk Classification Analysis

2.1 Prohibited Practices (Article 5)

Genesis does not engage in any prohibited practice under Article 5. The system does not deploy subliminal or manipulative techniques, exploit vulnerabilities of specific groups, process biometric data for categorization, perform social scoring,

assess criminal risk, compile facial recognition databases, infer emotions in workplaces or educational institutions, or perform real-time remote biometric identification. Its outreach capabilities could theoretically compose persuasive messages, but the system operates transparently to its operator and does not target vulnerable individuals. No further analysis is required under Article 5.

2.2 High-Risk Classification (Article 6 and Annex III)

Under Article 6, an AI system is high-risk if it either: (a) is used as a safety component or product covered by Annex I EU harmonisation legislation requiring third-party conformity assessment, or (b) falls under an Annex III use case category.

Annex I analysis: Genesis is not a product or safety component covered by any Annex I harmonisation legislation (medical devices, machinery, toys, marine equipment, etc.). Article 6(1) does not apply.

Annex III analysis: The Annex III categories are:

1. **Biometrics:** Not applicable. Genesis does not perform biometric identification, categorization, or emotion recognition.
2. **Critical infrastructure:** Not applicable. Genesis does not manage digital infrastructure, road traffic, or utility supply.
3. **Education and vocational training:** Not applicable. Genesis does not determine access to education, evaluate learning outcomes, or monitor students.
4. **Employment, workers management, and access to self-employment:** Not directly applicable. Genesis does not perform recruitment, evaluate candidates, or monitor worker performance in an employment context. However, if the system's outreach capabilities were used to contact hiring managers on behalf of its operator, this would be the operator using the system as a job application tool, not the system performing employment-related AI functions as contemplated by Annex III.
5. **Essential public and private services:** Not applicable. Genesis does not assess eligibility for benefits, evaluate creditworthiness, or classify emergency calls.
6. **Law enforcement:** Not applicable.
7. **Migration, asylum, and border control:** Not applicable.
8. **Administration of justice and democratic processes:** Not applicable.

Article 6(3) exception analysis: Even if a system falls under an Annex III category, it is not considered high-risk if it performs a narrow procedural task, improves the result of a previously completed human activity, detects decision-making patterns without replacing human assessment, or performs a preparatory task. Genesis's research and outreach capabilities would likely qualify for one or more of these exceptions (the system performs preparatory tasks and improves the result of human-initiated activities).

Conclusion: Genesis does not fall under any Annex III high-risk category. It is not classified as high-risk under Article 6.

2.3 Limited Risk-Transparency Obligations (Article 50)

Article 50 imposes transparency obligations on certain AI systems regardless of their risk classification.

Article 50(1)-Interaction transparency: Providers must design AI systems intended to interact directly with natural persons so that those persons are informed they are interacting with an AI system. This obligation applies unless it is obvious from the circumstances and context of use.

Application to Genesis: When Genesis sends messages to external parties (outreach emails, Telegram messages, forum posts), the recipients are interacting with an AI-generated or AI-assisted communication. They are natural persons. Under Article 50(1), these recipients should be informed that the communication involves an AI system.

Current implementation: Genesis does not systematically disclose its AI nature in outreach communications. The operator may or may not include such disclosure. This is a compliance gap.

Article 50(2)-Synthetic content marking: Providers of AI systems generating synthetic audio, image, video, or text content must mark that content as artificially generated in a machine-readable format.

Application to Genesis: When Genesis generates text content (outreach messages, forum posts, composed emails), that content is synthetic text. Under Article 50(2), this content should be marked as AI-generated.

Current implementation: No machine-readable marking is applied to AI-generated content. This is a compliance gap, though the implementing acts defining the specific marking requirements had not been finalized at the time of this assessment.

Article 50(4)-Deployer obligations: Deployers of AI systems that generate synthetic content must disclose that the content has been artificially generated or manipulated.

Application to Genesis: The operator, as deployer, has an obligation to disclose AI generation of synthetic content. This overlaps with Article 50(2) but places the obligation on the deployer rather than the provider.

Conclusion: Genesis falls under Limited Risk with transparency obligations under Article 50. The primary compliance gap is the absence of AI disclosure in outreach communications.

2.4 Risk Classification Summary

Classification	Applicable?	Basis
Prohibited	No	No prohibited practices implicated
High-Risk	No	No Annex III use case applies

Classification	Applicable?	Basis
Limited Risk (transparency)	Yes	Article 50-interaction and synthetic content disclosure
Minimal Risk	Partial	Non-outreach functions fall under minimal risk

3. Obligation Mapping

Given the Limited Risk classification, the primary obligations are transparency-related. However, for completeness and because the system's capabilities could change (e.g., expansion into use cases that trigger high-risk classification), this section also maps the high-risk obligations (Articles 9-17) to identify what would be required and what is already in place.

3.1 Current Obligations (Limited Risk)

Article 50(1)-Interaction disclosure

Requirement: Persons interacting with an AI system must be informed they are interacting with AI, unless obvious from context.

Current state: Not implemented for outreach communications.

Remediation: Include AI disclosure in outreach messages. This can be as simple as a footer line ("This message was composed with AI assistance") or as systematic as modifying the outreach pipeline to append disclosure automatically. The operator retains discretion over whether a particular communication warrants disclosure (some contexts may make AI involvement obvious).

Article 50(2)-Synthetic content marking

Requirement: AI-generated content must be marked as artificially generated in a machine-readable format.

Current state: Not implemented. No metadata marking exists.

Remediation: Pending publication of implementing acts defining marking standards. When standards are published, the outreach pipeline should apply appropriate metadata. For text content, this may involve C2PA (Coalition for Content Provenance and Authenticity) metadata or equivalent.

Article 50(4)-Deployer disclosure

Requirement: Deployers must disclose that content has been AI-generated.

Current state: Not systematically implemented. The operator may disclose voluntarily.

Remediation: Same as Article 50(1)–systematic disclosure in outreach communications.

3.2 Hypothetical High-Risk Analysis (Articles 9-17)

This section is included to demonstrate readiness, not because these obligations currently apply. If Genesis were classified as high-risk (for example, through future expansion into an Annex III use case), the following obligations would apply.

Article 9-Risk management system: Requires a continuous lifecycle risk management system. Partially met–action classification, approval gates, and the NIST assessment (D1) constitute elements of such a system, but no formal risk management document exists.

Article 10-Data governance: Requires training and testing data to be relevant, representative, and error-free. Partially applicable–Genesis does not train models, but its memory and knowledge stores function as behavioral conditioning data. Quality controls exist (memory verification, credibility scoring) but are not formalized.

Article 11-Technical documentation: Requires Annex IV-compliant documentation before market placement. Partially met–architecture docs and this assessment series exist, but a formal Annex IV package has not been prepared.

Article 12-Record-keeping: Requires automatic logging of risk-relevant events. Met–the event bus, approval records, session transcripts, and Git history provide structured audit trails. This is one of Genesis’s strongest compliance areas.

Article 13-Transparency for deployers: Requires the system to be understandable by deployers. Met for the current model (operator is deployer and builder). Distribution to third-party deployers would require additional documentation.

Article 14-Human oversight: Requires effective human oversight including the ability to intervene, interrupt, and override. Met–the autonomy model, approval gates, and kill switch provide layered oversight. This is a central design principle.

Article 15-Accuracy, robustness, and cybersecurity: Requires appropriate levels of each. Partially met–cybersecurity is addressed through defense-in-depth; accuracy is assessed via correction rates; robustness via circuit breakers and provider diversification. Formal metrics are not defined.

Article 17-Quality management system: Requires a formal QMS. Partially met–development processes (testing, code review, PR workflow) exist but are not formalized as a QMS.

3.3 Readiness Summary for High-Risk Obligations

Article	Requirement	Readiness
9	Risk management	Partial–elements exist, no formal system
10	Data governance	Partial–applicable controls exist for memory/knowledge data

Article	Requirement	Readiness
11	Technical documentation	Partial-exists but not Annex IV format
12	Record-keeping	Met-event bus, approval records, transcripts
13	Transparency	Met-for current operator-as-deployer model
14	Human oversight	Met-central design principle
15	Accuracy/robustness/security	Partial-security strong, metrics undefined
17	Quality management	Partial-processes exist, not formalized

3.4 Preparatory Remediation for Hypothetical High-Risk

Even though high-risk classification does not currently apply, three partially-met areas warrant preparatory attention. If classification changed, these would be the hardest to address under time pressure:

Article 9 (Risk management system): Formalize existing controls into a documented risk management system. The action classification taxonomy, approval gates, and correction mechanism already function as risk controls. The gap is documentation and lifecycle process, not capability.

Article 15 (Accuracy, robustness, cybersecurity): Define formal metrics for action classification accuracy and behavioral consistency. Current assessment is qualitative (correction rates, approval patterns). A high-risk designation would require quantitative accuracy metrics with defined thresholds.

Article 17 (Quality management system): Consolidate development processes (PR workflow, testing, code review, deployment procedures) into a documented QMS. The processes exist but are not formalized as a management system.

If high-risk classification were triggered, the primary work would be: formalizing existing processes into documented systems (risk management, quality management, technical documentation per Annex IV) rather than building new capabilities from scratch.

4. General-Purpose AI Model Obligations (Chapter V)

4.1 Genesis and GPAI

Genesis does not provide a GPAI model. It consumes GPAI models (Anthropic Claude, DeepInfra, Groq, Mistral) via API. The GPAI obligations fall on the model providers, not on Genesis as a downstream deployer.

However, the GPAI provisions are relevant to Genesis in two ways:

Provider obligations affect Genesis indirectly: GPAI model providers must provide technical documentation and instructions for use to downstream providers. Genesis depends on this documentation to understand model capabilities and limitations. If a provider's documentation is inadequate, Genesis's ability to make informed risk management decisions is impaired.

Systemic risk models: GPAI models present systemic risk when the cumulative compute used for training exceeds 10^{25} FLOPs. Anthropic Claude is trained at a scale that likely exceeds this threshold. Providers of systemic-risk models must conduct model evaluations, adversarial testing, serious incident tracking, and cybersecurity protections. Genesis benefits from these requirements but cannot verify compliance.

4.2 Downstream Deployer Considerations

As a deployer of GPAI-based AI systems, the Genesis operator should:

1. Monitor provider compliance disclosures (training data summaries, technical documentation, copyright compliance statements)
2. Assess whether provider documentation is sufficient for Genesis's risk management needs
3. Track provider model version changes and evaluate whether they affect Genesis's governance properties
4. Maintain awareness of GPAI codes of practice as they are published

Current state: None of these are systematically performed. The NIST assessment (D1, GOV-6) identified third-party governance as the most significant gap.

5. Territorial Scope

5.1 Applicability to Genesis

The EU AI Act applies to: - Providers placing AI systems on the EU market (regardless of location) - Deployers located in the EU - Providers and deployers located outside the EU where the AI system's output is used in the EU

Genesis is operated by a US-based individual. The system is not placed on the EU market as a product. However, if Genesis sends outreach messages to EU residents, generates content consumed by EU residents, or interacts with EU-based web services, the Act's output-use provisions may apply.

5.2 Practical Assessment

For the current deployment context (personal use, US-based operator), EU AI Act obligations are relevant but enforcement risk is low. The primary practical concern is Article 50 transparency obligations for outreach to EU residents.

If Genesis were distributed as an open-source product that EU-based users deploy, the provider obligations under the Act would apply more directly. The open-source

exemptions in the Act (free and open-license models are exempt from most GPAI obligations unless they present systemic risk) apply to GPAI models, not to AI systems. An open-source AI system used as a high-risk system would still be subject to high-risk requirements.

6. Timeline Analysis

The EU AI Act entered into force with staggered deadlines:

- **6 months** (February 2025): Prohibited AI practices took effect. Not applicable to Genesis.
- **12 months** (August 2025): GPAI obligations took effect. Applicable to Genesis's LLM providers, not directly to Genesis.
- **24 months** (August 2026): High-risk AI systems under Annex III. Not currently applicable to Genesis, but this deadline is approaching and would apply if classification changed.
- **36 months** (August 2027): High-risk AI systems under Annex I. Not applicable.

For Genesis, the relevant timeline is: - **Now**: Article 50 transparency obligations are in effect. The system should implement AI disclosure in outreach communications. - **August 2026**: If Genesis's use expanded into an Annex III category, high-risk obligations would apply from this date. - **Ongoing**: Monitor GPAI provider compliance and codes of practice as they are published.

7. The "Agent" Question

The EU AI Act does not define "agent" or "autonomous agent" as a legal category. The Act regulates "AI systems" defined as machine-based systems that generate outputs (predictions, recommendations, decisions, content) that can influence environments. Genesis fits this definition.

The Act's risk classification is use-case-based (Annex III categories), not capability-based. A system that can take autonomous external actions does not automatically receive a higher risk classification unless those actions fall within a regulated use case. This means the Act's risk framework has a structural gap for autonomous agents: a system that can send messages, execute code, browse the web, and interact with services is classified the same as a chatbot that answers questions, unless it performs a specific Annex III function.

This gap is not unique to the EU AI Act. It reflects a broader challenge in AI governance: existing frameworks regulate AI by what it does (classification, recommendation, decision support), not by how much authority it has. An autonomous agent that has broad authority to act but does not perform any specific Annex III function falls through the classification framework.

The NIST assessment (D1) and ISO 42001 analysis (D2) identify similar adaptation challenges. The methodology document (D4) discusses this pattern across all three frameworks.

8. Limitations of This Assessment

This assessment has inherent scope constraints that affect how the findings should be interpreted:

- **Single-system basis.** The risk classification and obligation mapping reflect one specific agent architecture. Systems with different capability profiles, deployment contexts, or operational patterns may classify differently under the same provisions.
 - **No legal counsel review.** This is an engineering-informed governance assessment, not a legal opinion. The classification analysis follows the text of the Regulation and published interpretive guidance, but specific determinations would require qualified EU legal review.
 - **Implementing acts not yet published.** Several obligations (particularly Article 50(2) synthetic content marking) depend on implementing acts and harmonized standards that were still in development at the time of this assessment. The specific requirements may differ from what the Regulation's text implies.
 - **Temporal snapshot.** The EU AI Act's staggered implementation means the regulatory environment is actively changing. High-risk system obligations take effect August 2026; compliance expectations may shift as guidance is published.
-

9. Gap Summary and Remediation

Current Compliance Gaps

Article 50(1)-Interaction disclosure: Not implemented. AI disclosure not included in outreach communications.

Article 50(2)-Synthetic content marking: Not implemented. No machine-readable AI generation metadata applied.

Article 50(4)-Deployer disclosure: Not systematically implemented.

Remediation Plan

Immediate (no technical changes required): 1. Include AI disclosure in outreach message templates ("This message was composed with AI assistance" or equivalent) 2. Add AI disclosure guidance to outreach pipeline documentation

Short-term (minor development): 3. Modify the outreach pipeline to automatically append AI disclosure to external communications, with operator override capability 4. Track AI disclosure compliance in the outreach delivery log

Medium-term (dependent on external standards): 5. Implement machine-readable synthetic content marking when implementing acts specify the required format (expected C2PA or equivalent) 6. Implement metadata preservation for AI-generated content that may be redistributed

Ongoing: 7. Monitor GPAI provider compliance disclosures 8. Re-evaluate risk classification if system capabilities or use patterns change (especially expansion into Annex III use case categories) 9. Track publication of harmonized standards and codes of practice

Cross-Reference to Other Assessments

The following gaps are identified in multiple assessments and should be addressed as a coordinated effort:

- **Third-party governance** (NIST GOV-6, ISO A.10, EU AI Act GPAI Section 4.2): Assess and document LLM provider data handling, terms of service, and compliance posture.
- **External impact detection** (NIST MANAGE-4.3, ISO A.8.3/A.8.4): Implement mechanisms for detecting and communicating about incidents affecting external parties.
- **AI disclosure in communications** (EU AI Act Article 50, ISO A.8.5): Systematic transparency about AI involvement in external communications.